

Semantic Spotter – Project Submission

1. Overview

This project demonstrates a **Retrieval-Augmented Generation (RAG)** system tailored for the **insurance domain**, built using the LangChain framework.

2. Problem Statement

The objective is to develop a **generative search system** that can accurately and efficiently answer questions using the contents of various policy documents.

3. Dataset

The insurance policy documents used in this project are available in the DocumentsPolicy directory.

4. Methodology

LangChain is a modular framework that simplifies the development of LLM-driven applications by offering reusable components, tools, and integrations.

Key Features:

- **Modular Components** – Prompt templates, chains, agents, memory, callbacks, etc.
 - **Composability** – Build pipelines by combining reusable blocks.
 - **LLM-Agnostic** – Supports various providers like OpenAI, Cohere, Hugging Face.
 - **External Integration** – Easily connects to APIs, databases, and document sources.
-

5. System Components

Document Loading

- Uses `PyPDFDirectoryLoader` to load PDFs from the specified folder.

Text Splitting

- Splits documents into chunks using `RecursiveCharacterTextSplitter`, which intelligently segments text to preserve context.

Embedding Generation

- Converts text to vectors using `OpenAIEmbeddings`, enabling similarity search and semantic comparison.

Vector Storage

- Stores embeddings in `ChromaDB`, backed by `CacheBackedEmbeddings` for performance and caching.

Document Retrieval

- Uses `VectorStoreRetriever` to fetch relevant chunks based on similarity with the user query.

Re-Ranking with Cross Encoder

- Applies `HuggingFaceCrossEncoder` (`BAAI/bge-reranker-base`) to re-rank retrieved chunks based on contextual relevance.

Chains

- Uses `LLMChain` and LangChain Hub's `rlm/rag-prompt` to format queries, combine inputs, and generate final responses from the LLM.

6. System Architecture

The system follows a modular, four-stage Retrieval-Augmented Generation (RAG) pipeline:

1. **Load**
 - Source documents are ingested from formats like PDFs, JSON, plain text, images, URLs, and HTML.
 - LangChain's document loaders (e.g., `PyPDFDirectoryLoader`) are used to handle structured and unstructured content.
2. **Split**
 - Large documents are divided into smaller, semantically coherent chunks using tools like `RecursiveCharacterTextSplitter`.
 - This ensures that each chunk fits within token limits and retains meaningful context.
3. **Embed**
 - Each chunk is transformed into a dense vector using an embedding model such as `OpenAIEmbeddings`.
 - The embedding step converts natural language into a numeric form suitable for similarity comparison.
4. **Store**

- Vectors are stored in a vector database like ChromaDB for fast semantic retrieval.
- Optionally, `CacheBackedEmbeddings` can be used to avoid redundant embedding computations.

At query time, the same embedding model converts the user query into a vector. The system performs a similarity search to retrieve the most relevant document chunks, optionally reranks them using a cross-encoder, and finally passes them to the LLM to generate a context-aware response.

7. Requirements

- Python 3.7 or higher
 - `langchain==0.3.13`
 - OpenAI API Key (store it in a file named `OpenAI_API_Key.txt`)
-

8. Getting Started

Clone the Repository

```
git clone https://github.com/jafarijason/semantic-spotter-project.git
cd semantic-spotter-project
```

Open the Notebook

Launch and run all cells in the following notebook: `semantic-spotter-langchain-notebook.ipynb`

Manual Setup Instructions

```
virtualenv venv
source venv/bin/activate
pip install -r requirements.txt

jupyter notebook \
    --notebook-dir="." \
    --ip=0.0.0.0 --port=3225
```

Contributors

Jason Jafari